

```
import pandas as pd
import numpy as np

data=pd.read_csv('/content/twitter_training.csv',header=None)
data
```

	0	1	2	3
0	2401	Borderlands	Positive	im getting on borderlands and i will murder yo...
1	2401	Borderlands	Positive	I am coming to the borders and I will kill you...
2	2401	Borderlands	Positive	im getting on borderlands and i will kill you ...
3	2401	Borderlands	Positive	im coming on borderlands and i will murder you...
4	2401	Borderlands	Positive	im getting on borderlands 2 and i will murder ...
...
74677	9200	Nvidia	Positive	Just realized that the Windows partition of my...
74678	9200	Nvidia	Positive	Just realized that my Mac window partition is ...
74679	9200	Nvidia	Positive	Just realized the windows partition of my Mac ...
74680	9200	Nvidia	Positive	Just realized between the windows partition of...
74681	9200	Nvidia	Positive	Just like the windows partition of my Mac is l...

74682 rows × 4 columns

```
data.columns=['ID','LOCATION','TARGET','TEXT']
data
```

	ID	LOCATION	TARGET	TEXT
0	2401	Borderlands	Positive	im getting on borderlands and i will murder yo...
1	2401	Borderlands	Positive	I am coming to the borders and I will kill you...
2	2401	Borderlands	Positive	im getting on borderlands and i will kill you ...
3	2401	Borderlands	Positive	im coming on borderlands and i will murder you...
4	2401	Borderlands	Positive	im getting on borderlands 2 and i will murder ...
...
74677	9200	Nvidia	Positive	Just realized that the Windows partition of my...
74678	9200	Nvidia	Positive	Just realized that my Mac window partition is ...
74679	9200	Nvidia	Positive	Just realized the windows partition of my Mac ...
74680	9200	Nvidia	Positive	Just realized between the windows partition of...
74681	9200	Nvidia	Positive	Just like the windows partition of my Mac is l...

74682 rows × 4 columns

```
data.info()

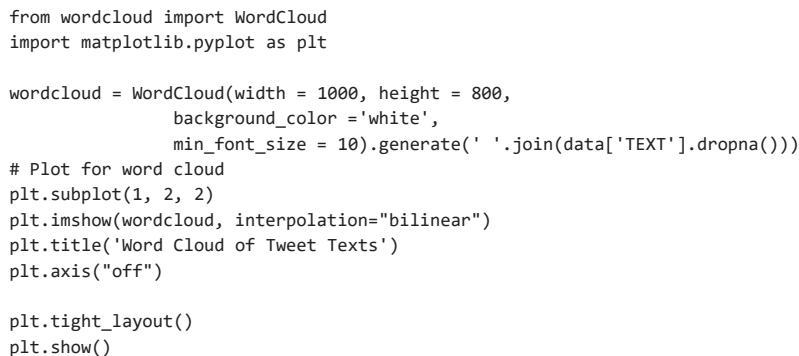
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 74682 entries, 0 to 74681
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0    ID           74682 non-null  int64
1    LOCATION     74682 non-null  object
2    TARGET       74682 non-null  object
3    TEXT         73996 non-null  object
dtypes: int64(1), object(3)
memory usage: 2.3+ MB

data.isna().sum()

ID           0
LOCATION       0
TARGET       0
TEXT        686
dtype: int64

data['TARGET'].value_counts()
```

```
<Axes: xlabel='TARGET', ylabel='count'>
```



	ID	LOCATION	TARGET	TEXT
0	2401	Borderlands	Positive	im getting on borderlands and i will murder yo...
1	2401	Borderlands	Positive	I am coming to the borders and I will kill you...
2	2401	Borderlands	Positive	im getting on borderlands and i will kill you ...
3	2401	Borderlands	Positive	im coming on borderlands and i will murder you...
4	2401	Borderlands	Positive	im getting on borderlands 2 and i will murder ...
...
74677	9200	Nvidia	Positive	Just realized that the Windows partition of my...
74678	9200	Nvidia	Positive	Just realized that mv Mac window partition is ...

```
data.drop(['ID', 'LOCATION'],axis=1,inplace=True)
data
```

	TARGET	TEXT
0	Positive	im getting on borderlands and i will murder yo...
1	Positive	I am coming to the borders and I will kill you...
2	Positive	im getting on borderlands and i will kill you ...
3	Positive	im coming on borderlands and i will murder you...
4	Positive	im getting on borderlands 2 and i will murder ...
...
74677	Positive	Just realized that the Windows partition of my...
74678	Positive	Just realized that my Mac window partition is ...
74679	Positive	Just realized the windows partition of my Mac ...
74680	Positive	Just realized between the windows partition of...
74681	Positive	Just like the windows partition of my Mac is l...

61692 rows × 2 columns

```
data['TARGET']=data['TARGET'].map({'Positive':1,'Negative':-1,'Neutral':0})
data
```

	TARGET	TEXT
0	1	im getting on borderlands and i will murder yo...
1	1	I am coming to the borders and I will kill you...
2	1	im getting on borderlands and i will kill you ...
3	1	im coming on borderlands and i will murder you...
4	1	im getting on borderlands 2 and i will murder ...
...
74677	1	Just realized that the Windows partition of my...
74678	1	Just realized that my Mac window partition is ...
74679	1	Just realized the windows partition of my Mac ...
74680	1	Just realized between the windows partition of...
74681	1	Just like the windows partition of my Mac is l...

61692 rows × 2 columns

```
tweets=data.TEXT
tweets
```

```
0      im getting on borderlands and i will murder yo...
1      I am coming to the borders and I will kill you...
2      im getting on borderlands and i will kill you ...
3      im coming on borderlands and i will murder you...
4      im getting on borderlands 2 and i will murder ...
...
74677  Just realized that the Windows partition of my...
74678  Just realized that my Mac window partition is ...
74679  Just realized the windows partition of my Mac ...
74680  Just realized between the windows partition of...
74681  Just like the windows partition of my Mac is l...
Name: TEXT, Length: 61692, dtype: object
```

```
tweets=tweets.str.replace('[^a-zA-Z0-9]+',' ')
tweets
```

```
<ipython-input-14-243a49c37bfd>:1: FutureWarning: The default value of regex will change from True to False in a future version.
  tweets=tweets.str.replace('[^a-zA-Z0-9]+',' ')
0      im getting on borderlands and i will murder yo...
1      I am coming to the borders and I will kill you...
2      im getting on borderlands and i will kill you ...
3      im coming on borderlands and i will murder you...
4      im getting on borderlands 2 and i will murder ...
...
74677   Just realized that the Windows partition of my...
74678   Just realized that my Mac window partition is ...
74679   Just realized the windows partition of my Mac ...
74680   Just realized between the windows partition of...
74681   Just like the windows partition of my Mac is l...
Name: TEXT, Length: 61692, dtype: object
```

```
from nltk.stem import SnowballStemmer
from nltk.tokenize import TweetTokenizer
stemmer = SnowballStemmer('english')
tokenizer = TweetTokenizer()
tweets = tweets.fillna('').astype(str)
tweets = tweets.apply(lambda x: ' '.join([stemmer.stem(word.lower()) for word in tokenizer.tokenize(x)]))
print(tweets)
```

```
0      im get on borderland and i will murder you all
1      i am come to the border and i will kill you all
2      im get on borderland and i will kill you all
3      im come on borderland and i will murder you all
4      im get on borderland 2 and i will murder you m...
...
74677   just realiz that the window partit of my mac i...
74678   just realiz that my mac window partit is 6 yea...
74679   just realiz the window partit of my mac is now...
74680   just realiz between the window partit of my ma...
74681   just like the window partit of my mac is like ...
Name: TEXT, Length: 61692, dtype: object
```

```
from nltk.corpus import stopwords
import nltk
nltk.download('stopwords')
stop=stopwords.words('english')
tweets=tweets.apply(lambda x:[i for i in tk.tokenize(x) if i not in stop]).apply(lambda x:' '.join(x))
tweets
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
0      im get borderland murder
1      come border kill
2      im get borderland kill
3      im come borderland murder
4      im get borderland 2 murder
...
74677   realiz window partit mac like 6 year behind nv...
74678   realiz mac window partit 6 year behind nvidia ...
74679   realiz window partit mac 6 year behind nvidia ...
74680   realiz window partit mac like 6 year behind nv...
74681   like window partit mac like 6 year behind driv...
Name: TEXT, Length: 61692, dtype: object
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
vec=TfidfVectorizer()
train_data_vec=vec.fit_transform(tweets)
print(train_data_vec)
```

```
(0, 12683)    0.6730103007248203
(0, 4007)     0.42392218796406755
(0, 8437)     0.32424437145094054
(0, 9864)     0.5120670866716872
(1, 10846)    0.4942247167249265
(1, 4005)     0.7462672532242293
(1, 5138)     0.4459003432868494
(2, 10846)    0.5368353840071853
(2, 4007)     0.4835608885775419
(2, 8437)     0.3698600847671917
(2, 9864)     0.5841062875983802
(3, 5138)     0.40949932608130163
(3, 12683)    0.6490608926334446
(3, 4007)     0.40883670492226853
(3, 9864)     0.493844923332353
(4, 12683)    0.6730103007248203
(4, 4007)     0.42392218796406755
(4, 8437)     0.32424437145094054
```

```

(4, 9864)      0.5120670866716872
(5, 12683)     0.6730103007248203
(5, 4007)      0.42392218796406755
(5, 8437)      0.32424437145094054
(5, 9864)      0.5120670866716872
(6, 1628)      0.2956243164240888
(6, 12415)     0.2956243164240888
:             :
(61689, 15162) 0.3220233159223399
(61689, 20415) 0.20408119305150327
(61690, 13875) 0.41830700432991136
(61690, 13261) 0.19568691019259787
(61690, 6577)  0.2840759041157257
(61690, 19974) 0.26694675536221796
(61690, 4476)  0.2786242145665763
(61690, 11727) 0.3284995782133617
(61690, 13171) 0.2971116435226657
(61690, 3556)  0.28772771983646384
(61690, 9785)  0.25316576582403766
(61690, 15162) 0.2931579689778407
(61690, 7205)  0.20429811615316418
(61690, 8094)  0.1652258772383998
(61690, 20415) 0.1857878765399137
(61690, 11353) 0.15675220221538724
(61691, 13875) 0.46686223079192973
(61691, 6577)  0.3170501783066094
(61691, 19974) 0.2979327537455749
(61691, 11727) 0.3666303559620592
(61691, 13171) 0.33159904867347145
(61691, 3556)  0.3211258806404872
(61691, 9785)  0.2825521279092433
(61691, 20415) 0.20735331131847726
(61691, 11353) 0.34989460874581124

```

```
train_data_vec.shape
```

```
(61692, 20705)
```

```
y=data['TARGET'].values
```

```
y
```

```
array([1, 1, 1, ..., 1, 1, 1])
```

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(train_data_vec,y,test_size=0.2,random_state=1)
```

```
from sklearn import svm
```

```
from sklearn.naive_bayes import MultinomialNB
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.ensemble import AdaBoostClassifier
```

```
svm_model=svm.SVC()
```

```
nb_model=MultinomialNB()
```

```
rf_model=RandomForestClassifier()
```

```
ab_model=AdaBoostClassifier()
```

```
lstmodel=[svm_model,nb_model,rf_model,ab_model]
```

```
from sklearn.metrics import ConfusionMatrixDisplay,classification_report
```

```
for i in lstmodel:
```

```
    print(i)
```

```
    i.fit(x_train,y_train)
```

```
    y_pred=i.predict(x_test)
```

```
    print('*****')
```

```
    print(classification_report(y_test,y_pred))
```

```
    print('*****')
```

```
    print(ConfusionMatrixDisplay.from_predictions(y_test,y_pred))
```

```

SVC()
*****
      precision    recall  f1-score   support

     -1         1.00      1.00      1.00         1
       1         1.00      1.00      1.00         1

 accuracy          1.00          1.00          1.00         2
 macro avg         1.00      1.00      1.00         2
 weighted avg      1.00      1.00      1.00         2

*****
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7ff7ab08c5e0>
MultinomialNB()
*****
      precision    recall  f1-score   support

     -1         1.00      1.00      1.00         1
       1         1.00      1.00      1.00         1

 accuracy          1.00          1.00          1.00         2
 macro avg         1.00      1.00      1.00         2
 weighted avg      1.00      1.00      1.00         2

*****
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7ff7ab08c5e0>
RandomForestClassifier()
*****
      precision    recall  f1-score   support

     -1         1.00      1.00      1.00         1
       1         1.00      1.00      1.00         1

 accuracy          1.00          1.00          1.00         2
 macro avg         1.00      1.00      1.00         2
 weighted avg      1.00      1.00      1.00         2

*****
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7ff7dd54c400>
AdaBoostClassifier()
*****
      precision    recall  f1-score   support

     -1         0.50      1.00      0.67         1
       1         0.00      0.00      0.00         1

 accuracy          0.50          0.50          0.50         2
 macro avg         0.25      0.50      0.33         2
 weighted avg      0.25      0.50      0.33         2

*****
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay object at 0x7ff7dd54c400>
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are
_warn_prf(average, modifier, msg_start, len(result))

```

